

Part A: Theoretical Concepts

1. Activation Functions

Define and compare the following activation functions: Sigmoid, ReLU, Tanh, and Leaky ReLU.

(a) *Sigmoid*:

- **Formula:** $\sigma(x) = \frac{1}{1 + e^{-x}}$
- **Range:** $(0, 1)$
- **Use Case:** Used in binary classification tasks.
- **Limitation:** Vanishing gradients during backpropagation for large or small values of x .

(b) *ReLU (Rectified Linear Unit)*:

- **Formula:** $f(x) = \max(0, x)$
- **Range:** $[0, \infty)$
- **Use Case:** Default activation for deep networks.
- **Limitation:** Dead neurons (outputs zero for negative inputs).

(c) *Tanh*:

- **Formula:** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- **Range:** $(-1, 1)$
- **Use Case:** Normalized data-centered activations.
- **Limitation:** Vanishing gradients.

(d) *Leaky ReLU*:

- **Formula:** $f(x) = \max(0.01x, x)$
- **Range:** $(-\infty, \infty)$
- **Use Case:** Solves the dead neurons problem.
- **Limitation:** Requires manual tuning of the slope for negative inputs.

2. Discussion of Optimization Algorithms

- **Comparison:**
 - **SGD (Stochastic Gradient Descent):**
 - Simple, faster for large datasets.
 - Sensitive to learning rate; may converge slowly.
 - **Adam (Adaptive Moment Estimation):**
 - Combines momentum and adaptive learning rates.
 - Suitable for sparse gradients.
 - **RMSprop (Root Mean Square Propagation):**
 - Adjusts learning rate using recent gradient magnitudes.
 - Performs well in RNNs and non-stationary problems.
- **Learning Rate Impact:**
 - High learning rate: Fast but may overshoot the minimum.
 - Low learning rate: Stable but slow convergence.
 - Modern optimizers adapt the learning rate dynamically.

Part B: Practical Implementation

1. Data Preprocessing :Download and preprocess the CIFAR-10 dataset

```
import tensorflow as tf

from tensorflow.keras.datasets import cifar10

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

x_train, x_test = x_train / 255.0, x_test / 255.0

data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomRotation(0.1)])
```

2. Model Design

- **Design a CNN:**
 - At least 3 convolutional layers and 2 fully connected layers.
 - Include regularization

Code:

```
from tensorflow.keras import models, layers

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(10, activation='softmax')])
```

3. Model Training

- **Compile and train the model:** Use early stopping or learning rate scheduling if necessary.

Code:

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)
```

```
history = model.fit( x_train, y_train, epochs=30,  
                    validation_data=(x_test, y_test),  
                    callbacks=[early_stopping])
```

4. Model Evaluation

- Evaluate the model on the test set.
- Generate a confusion matrix.

Code:

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
  
test_loss, test_acc = model.evaluate(x_test, y_test)  
  
print(f"Test Accuracy: {test_acc}")  
  
y_pred = model.predict(x_test).argmax(axis=1)  
  
y_true = y_test.argmax(axis=1)  
  
cm = confusion_matrix(y_true, y_pred)  
  
ConfusionMatrixDisplay(cm).plot()  
  
plt.show()
```

5. Error Analysis and conclusion

- Identify errors using the confusion matrix.
- **Example of errors:**
 1. Class A misclassified as Class B .
 2. Poor performance on smaller objects.
 3. Misclassification in overlapping classes.

Proposed Solutions for error:

- Increase training data diversity using augmentation.
- Use a pre-trained model for transfer learning.
- Fine-tune hyperparameters or increase model complexity.